

# Probabilistic Forecasting with Randomized/Quasi-Randomized networks

T. Moudiki

2024-07-03

# Plan

- ▶ 1 - Context
- ▶ 2 - Quasi-randomized *neural* networks (QRNs)
- ▶ 3 - QRN forecasting with `nnetsauce`

# Plan

- ▶ **1 - Context**
- ▶ 2 - Quasi-randomized *neural* networks (QRNs)
- ▶ 3 - QRN forecasting with `netsauce`

# 1 - Context

## ► Time series:

*A set of data collected sequentially usually at fixed intervals of time*  
(Merriam-Webster dictionary online)

## ► Objective:

- Forecasting one time series (**univariate forecasting**). **Example:** **Measurements of the annual flow of the river Nile** (in  $10^8 m^3$ ) at Aswan from 1871 to 1970 => What happens in 1971, 1972, etc.? How **“certain”** can we be about the forecast?
- Forecasting many related time series (**multivariate forecasting**). **Example:** Growth rates of personal **consumption** and personal **income** in the USA from 1970 to 2010 => What happens in 2011, 2012, etc.? How **“certain”** can we be?
- Forecasting tool? **Quasi-randomized *neural networks*** applied to time series lags
- Implemented in Python package `nnetsauce` version 0.22.4

# 1 - Context

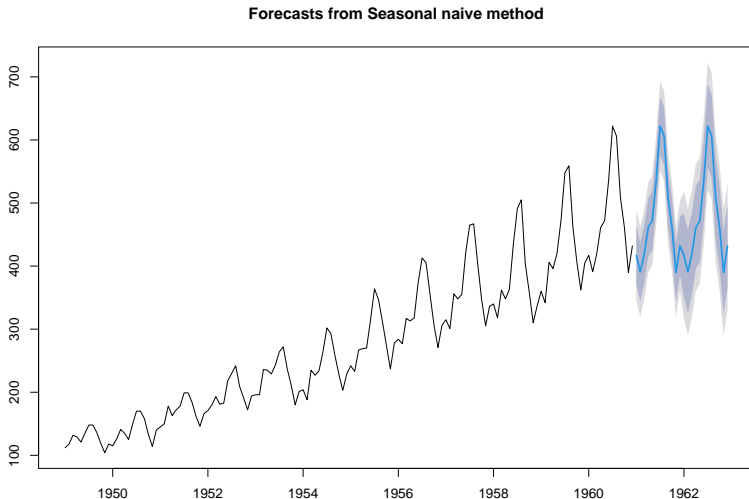
## Point forecasts/Uncertainty quantification

- ▶ **Point forecasts** for one time series and for multivariate time series: **answers** what happens in 1971, 1972, etc.? What happens in 2011, 2012, etc.?
- ▶ **Point forecasts**: cool, but not very informative. How wrong can we be, based on the assumptions that we made: **answers** how “*certain*” can we be about the forecast?
- ▶ **Uncertainty quantification** needed: **prediction intervals** and/or **predictive simulations**.
- ▶ **prediction intervals**: point forecast  $\pm$  a term (with a level of confidence)
- ▶ **predictive simulations**: future scenarios for the variables of interest

# 1 - Context

## Point forecasts/Uncertainty quantification

```
plot(forecast::snaive(AirPassengers))
```



# 1 - Context

## Uncertainty quantification for time series

- ▶ Based on:
  - ▶ **Bayesian priors**
  - ▶ **In-sample residuals** = model fit - true observation on the *whole training set*
  - ▶ **Calibrated residuals** = model fit - true observation on a *held-out calibration set*
- ▶ **Calibrated residuals** used in `nnetsauce` for methods based on **sequential split conformal prediction** (more on this later)

# 1 - Context

## Uncertainty quantification for time series



- ▶ Time series **residuals** matter:
  - ▶ can contain information  $\Rightarrow$  need to augment model capacity for example, so that they become smaller (and hopefully centered, with constant variance).
  - ▶ used for uncertainty quantification  $\Rightarrow$  need to infer their distribution, either parametric or by simulation.
  - ▶ assumption for residuals-based uncertainty quantification: **“pattern” observed in residuals is continuing** in the future.
- ▶ Important property of residuals: **stationarity**
  - ▶ **Stationarity**: no significant correlation of time series lags
  - ▶ **Stationarity**  $\Rightarrow$  time series mean and variance do not change over time.



# 1 - Context

## Uncertainty quantification for time series



### ► In-sample residuals



# 1 - Context

## Uncertainty quantification for time series

### ► Calibrated residuals?/Sequential split conformal prediction



# 1 - Context

## Uncertainty quantification for time series

- ▶ In `nnetsauce` version 0.22.4 (10 methods):
  - ▶ Can be achieved by a **Bayesian base learner**
  - ▶ Can be achieved by a **conformalized based learner**
  - ▶ Can rely on **in-sample** residuals for methods based on:
    - ▶ parametric residuals distribution inference (`gaussian`)
    - ▶ density estimation and simulation of residuals (`kde`)
    - ▶ bootstrap resampling (`bootstrap` and `block-bootstrap`)
  - ▶ Can rely on **calibrated** residuals for methods based on:
    - ▶ **sequential split conformal prediction** (`scp*-kde`, `scp*-bootstrap`, `scp*-block-bootstrap`)

# Plan

- ▶ 1 - Context
- ▶ 2 - **Quasi-randomized *neural* networks (QRNs)**
- ▶ 3 - QRN forecasting with `nnet`sauce

## 2 - Quasi-randomized *neural* networks (QRNs)

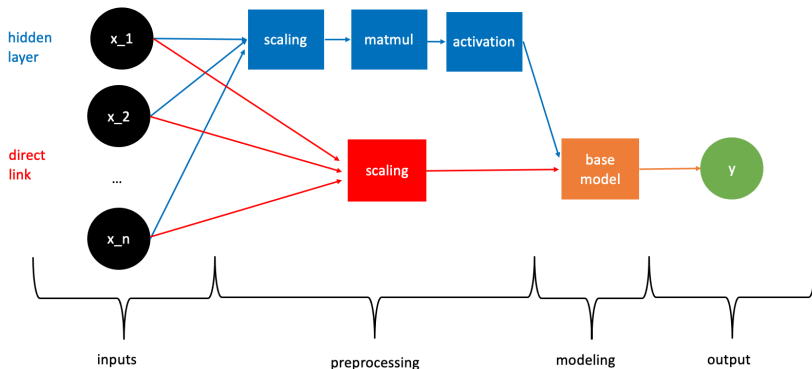


Figure 1: QRN principle

- ▶ A simple case (next slide): base model = **base learner** = Linear Regression

## 2 - Quasi-randomized *neural* networks (QRNs)

**Simple case:** **base learner** = Linear Regression;  $y \in \mathbb{R}^n$ , to be explained by  $X^{(j)}, j \in \{1, \dots, p\}$

$$y = \beta_0 + \sum_{j=1}^p \beta_j x^{(j)} + \sum_{l=1}^L \gamma_l g \left( \sum_{j=1}^p w^{(j,l)} x^{(j)} \right) + \epsilon$$

With:

- ▶  $g$  : **activation function**  $\rightarrow$  nonlinearity
- ▶  $L$  : number of nodes in the **hidden layer**
- ▶  $w^{(j,l)}$ , hidden layer: **pseudo/quasi-random**
- ▶ **Quasi-random**: designed to cover the space *parsimoniously*
- ▶  $\beta_j$  and  $\gamma_l$  : linear model coefficients
- ▶  $\epsilon$  : residuals

## 2 - Quasi-randomized *neural* networks (QRNs)

### QRNs applied to time series

- ▶ **Response**  $y$  = most recent time series observations
- ▶ **Covariates**  $X$  = time series lags
- ▶ **base learner**: can be any Machine Learning model
- ▶ **Multivariate forecasting case**: the base learner **shared** by all the time series

## 2 - Quasi-randomized *neural* networks (QRNs)

### QRNs applied to time series (illustration)

- ▶  $p = 2$  time series  $X^{(1)}$  and  $X^{(2)}$ ;  $n = 5$  dates  $t_1 < \dots < t_5$
- ▶  $X^{(2)} := (X_{t_1}^{(1)}, \dots, X_{t_5}^{(1)})$  and  $X^{(2)} := (X_{t_1}^{(2)}, \dots, X_{t_5}^{(2)})$
- ▶  $k = 2$  lags;  $L = 3$  nodes in hidden layer.

Response variables:

$$\mathbf{Y} = \begin{pmatrix} X_{t_5}^{(1)} & X_{t_5}^{(2)} \\ X_{t_4}^{(1)} & X_{t_4}^{(2)} \\ X_{t_3}^{(1)} & X_{t_3}^{(2)} \end{pmatrix}$$

Covariates:

$$\mathbf{X} = \begin{pmatrix} X_{t_4}^{(1)} & X_{t_3}^{(1)} & X_{t_4}^{(2)} & X_{t_3}^{(2)} \\ X_{t_3}^{(1)} & X_{t_2}^{(1)} & X_{t_3}^{(2)} & X_{t_2}^{(2)} \\ X_{t_2}^{(1)} & X_{t_1}^{(1)} & X_{t_2}^{(2)} & X_{t_1}^{(2)} \end{pmatrix}$$



## 2 - Quasi-randomized *neural* networks (QRNs)

### QRNs applied to time series (illustration)

- ▶  $p = 2$  time series  $X^{(1)}$  and  $X^{(2)}$
- ▶  $n = 5$  dates  $t_1 < \dots < t_5$ ,
- ▶  $X^{(2)} := (X_{t_1}^{(1)}, \dots, X_{t_5}^{(1)})$  and  $X^{(2)} := (X_{t_1}^{(2)}, \dots, X_{t_5}^{(2)})$
- ▶  $k = 2$  lags
- ▶  $L = 3$  nodes in hidden layer.

Coefficients in hidden layer:

$$\mathbf{W} = \begin{pmatrix} W^{(1,1)} & W^{(1,2)} & W^{(1,3)} \\ W^{(2,1)} & W^{(2,2)} & W^{(2,3)} \\ W^{(3,1)} & W^{(3,2)} & W^{(3,3)} \\ W^{(4,1)} & W^{(4,2)} & W^{(4,3)} \end{pmatrix}$$

# Plan

- ▶ 1 - Context
- ▶ 2 - Quasi-randomized *neural* networks (QRNs)
- ▶ 3 - **QRN forecasting with `netsauce`**

# Plan

- ▶ **3 - QRN forecasting with nnetsauce**
- ▶ 3 - 1 nnetsauce's description (Python version)
- ▶ 3 - 2 Install+import Python packages (including nnetsauce)
- ▶ 3 - 3 Import data for the demo
- ▶ 3 - 4 Using the `fit + predict` interface
- ▶ 3 - 5 Using GPUs
- ▶ 3 - 6 AutoML with LazyMTS
- ▶ 3 - 7 Time series cross-validation

## 3 - QRN forecasting with nnetsauce

### 3 - 1 nnetsauce's description (Python version)

- ▶ General-purpose tool for Machine Learning using **Randomized** and **Quasi-Randomized *neural* networks**
  - ▶ GitHub: <https://github.com/Techtonique/nnetsauce>
  - ▶ PyPI: <https://pypi.org/project/nnetsauce/>
  - ▶ Conda: <https://anaconda.org/conda-forge/nnetsauce>
- ▶ **Tasks:**
  - ▶ Classification
  - ▶ Regression
  - ▶ **Univariate/Multivariate time series forecasting**

### 3 - 1 nnettsauce's description (Python version) (cont'd)

- ▶ **Simple interface** for each model:
  - ▶ `fit`: fitting model to training data
  - ▶ `predict`: model inference on unseen data
- ▶ **GPU** version optimizes matrices multiplication using **JAX** (not *magical*)
- ▶ Classes MTS and DeepMTS for **time series forecasting**
- ▶ Here: **focus on MTS**
- ▶ Also: automated Machine Learning (**AutoML**) with class LazyMTS

### 3 - 2 Install+import Python packages (including nnetsauce)

```
pip install nnetsauce
pip install git+https://github.com/Techtonique/mlsauce.git
pip install xgboost
```

```
import nnetsauce as ns
import mlsauce as ms
import numpy as np
import pandas as pd
import seaborn as sns

from sklearn.linear_model import Ridge
from statsmodels.tsa.seasonal import STL


sns.set_theme(style="darkgrid")
```

### 3 - 3 Import data for the demo

**Univariate:** the famous AirPassengers data set: **Monthly Airline Passenger Numbers 1949-1960**

```
url = "https://raw.githubusercontent.com/Techtonique/"  
url += "datasets/main/time_series/univariate/"  
url += "a10.csv"  
df_a10 = pd.read_csv(url)  
df_a10.index = pd.DatetimeIndex(df_a10.date)  
df_a10.drop(columns=['date'], inplace=True)
```

```
df_a10.plot()
```






### 3 - 3 Import data for the demo (cont'd)

**Multivariate: Heater vs Ice cream sales** data set

```
url = "https://raw.githubusercontent.com/Techtonique/"
url += "datasets/main/time_series/multivariate/"
url += "ice_cream_vs_heater.csv"
df_temp = pd.read_csv(url)
df_temp.index = pd.DatetimeIndex(df_temp.date)
df_icecream = df_temp.drop(columns=['date']).diff() \
.dropna()
```

```
df_icecream.plot()
```



### 3 - 4 Using the fit + predict interface

- ▶ Gaussian
- ▶ Bayesian (Gaussian prior)
- ▶ Kernel Density Estimation (KDE)
- ▶ Conformalized base learner: TweedieRegressor + sequential split conformal

## 3 - 4 Using the fit + predict interface


### ► Gaussian

```
from sklearn.ensemble import BaggingRegressor

regr = ns.MTS(obj=BaggingRegressor(),
              type_pi="gaussian",
              lags=20,
              show_progress=False)
regr.fit(df_icecream);
regr.predict(h=30);
```

### 3 - 4 Using the fit + predict interface

```
regr.plot("heater", type_plot="pi")
```



### 3 - 4 Using the fit + predict interface


#### ► Bayesian (Gaussian prior)

```
from sklearn.linear_model import BayesianRidge

regr = ns.MTS(obj=BayesianRidge(),
              lags=15,
              show_progress=False)
regr.fit(df_a10);
regr.predict(h=40, return_std=True);
```

### 3 - 4 Using the fit + predict interface

```
regr.plot(type_plot="pi")
```



### 3 - 4 Using the fit + predict interface

#### ► SCP-KDE

```
from sklearn.linear_model import LassoLarsIC
```

```
regr = ns.MTS(obj=LassoLarsIC(),  
              type_pi="scp-kde",  
              replications=250,  
              kernel='gaussian',  
              lags=15,  
              show_progress=False)
```

```
regr.fit(df_a10);
```

```
regr.predict(h=40);
```


```
##      0%|          | 0/250 [00:00<?, ?it/s] 100%|#####|
```

```
##      0%|          | 0/250 [00:00<?, ?it/s] 100%|#####|
```



### 3 - 4 Using the fit + predict interface

```
regr.plot(type_plot="spaghetti")
```



### 3 - 4 Using the fit + predict interface

#### ► Conformalized base learner


```
from sklearn.linear_model import TweedieRegressor

obj0 = ns.PredictionInterval(obj=TweedieRegressor(),
                             method="splitconformal",
                             type_split="sequential",
                             level=95)

regr = ns.MTS(obj=obj0,
              lags=20,
              show_progress=False)
regr.fit(df_icecream);
regr.predict(h=30, return_pi=True);
```

### 3 - 4 Using the fit + predict interface

```
regr.plot("heater", type_plot="pi")
```



### 3 - 5 Using GPUs

- ▶ Public notebook on GitHub (<https://bit.ly/45RchgD>)
- ▶ Simulated multivariate time series: 100 series, 10000 observations
- ▶ Ran on **Kaggle notebooks**, with **accelerator GPU P100**

```
import numpy as np
import pandas as pd
import nnetsauce as ns
from sklearn.linear_model import Ridge
from time import time
```

```
df = generate_synthetic_mts(n_steps=10000, n_series=100,
                           amplitude=40, seed=14531)
df_ = df.diff().dropna()
```

### 3 - 5 Using GPUs (cont'd)

- **Example 1** on CPU, using Ridge regression

```
regr = Ridge()
obj_MTS = ns.MTS(regr,
                 lags = 15,
                 n_hidden_features=5,
                 nodes_sim="uniform",
                 backend="cpu", # specify backend
                 verbose = 1)

start = time()
obj_MTS.fit(df_)
print(f"Elapsed: {time()-start}")
```

Elapsed: 64.46652388572693

### 3 - 5 Using GPUs (cont'd)

- **Example 2** on GPU (uses **JAX** behind the scenes), using Ridge regression

```
regr = Ridge()
obj_MTS = ns.MTS(regr,
                 lags = 15,
                 n_hidden_features=5,
                 nodes_sim="uniform",
                 backend="gpu", # specify backend
                 verbose = 1)

start = time()
obj_MTS.fit(df_)
print(f"Elapsed: {time()-start}")
```

Elapsed: 40.53069853782654

- **37% time gain**

### 3 - 5 Using GPUs (cont'd)

- **Example 3** on GPU (uses **JAX** behind the scenes), using Ridge regression on GPU (mlsauce implementation) **too**

```
regr = ms.RidgeRegressor(reg_lambda=1.0, backend="gpu")
obj_MTS = ns.MTS(regr,
                  lags = 15,
                  n_hidden_features=5,
                  nodes_sim="uniform",
                  backend="gpu", # specify backend
                  verbose = 1)

start = time()
obj_MTS.fit(df_)
print(f"Elapsed: {time()-start}")
```

Elapsed: 23.551459312438965

- **63% time gain**

### 3 - 5 Using GPUs (cont'd)

#### ► Example 4 Using XGBoost

```
import xgboost as xgb

xgb1 = xgb.XGBRegressor()
xgb2 = xgb.XGBRegressor(tree_method = "gpu_hist")

obj_MTS = ns.MTS(xgb1,
                 lags = 15,
                 n_hidden_features=5,
                 nodes_sim="uniform",
                 backend="cpu", # specify backend
                 verbose = 1)

start = time()
obj_MTS.fit(df_)
print(f"Elapsed: {time()-start}")
```

Elapsed: 3285.469851732254



### 3 - 5 Using GPUs (cont'd)

► **Example 5** Using XGBoost on GPU

```
obj_MTS = ns.MTS(xgb1,
                  lags = 15,
                  n_hidden_features=5,
                  nodes_sim="uniform",
                  backend="gpu", # specify backend
                  verbose = 1)

start = time()
obj_MTS.fit(df_)
print(f"Elapsed: {time()-start}")
```

Elapsed: 3266.556359767914

► No substantial gain (19 seconds), **but...**

### 3 - 5 Using GPUs (cont'd)

► **Example 2** Using XGBoost's GPU implementation

```
obj_MTS = ns.MTS(xgb2,
                 lags = 15,
                 n_hidden_features=5,
                 nodes_sim="uniform",
                 backend="cpu", # specify backend
                 verbose = 1)

start = time()
obj_MTS.fit(df_)
print(f"Elapsed: {time()-start}")
```

Elapsed: 249.40361714363098

► **92% time gain** over backend="cpu"

### 3 - 5 Using GPUs (cont'd)

#### ► **Example 2** Using XGBoost's GPU implementation

```
obj_MTS = ns.MTS(xgb2,  
                 lags = 15,  
                 n_hidden_features=5,  
                 nodes_sim="uniform",  
                 backend="gpu", # specify backend  
                 verbose = 1)  
  
start = time()  
obj_MTS.fit(df_)  
print(f"Elapsed: {time()-start}")
```

Elapsed: 249.10436272621155

- a little faster than backend="cpu"
- **92% time gain** over backend="gpu" with XGBoost CPU

**Conclusion:** No magic trick, experiment to see what works *best*!

### 3 - 6 Automated Machine Learning (AutoML) with LazyMTS

```
# split data into training/testing set
idx_train = int(df_icecream.shape[0]*0.9)
idx_end = df_icecream.shape[0]
df_train = df_icecream.iloc[0:idx_train,:]
df_test = df_icecream.iloc[idx_train:idx_end,:]


# Train + predict on 5 ML models
regr = ns.LazyMTS(verbose=1, ignore_warnings=False,
                  custom_metric=None, show_progress=False,
                  preprocess=True,
                  type_pi="kde",
                  kernel="gaussian",
                  lags = 20, replications=250,
                  estimators=["Ridge",
                              "ElasticNet", "ExtraTreesRegressor",
                              "RandomForestRegressor"]);
models, predictions = regr.fit(df_train, df_test);
```

### 3 - 6 Automated Machine Learning (AutoML) with LazyMTS (cont'd)

```
print(models[['RMSE', 'WINKLERSCORE']])
```

| ##                            | RMSE | WINKLERSCORE |
|-------------------------------|------|--------------|
| ## Model                      |      |              |
| ## MTS(Ridge)                 | 8.52 | 220.35       |
| ## MTS(RandomForestRegressor) | 8.79 | 239.63       |
| ## MTS(ExtraTreesRegressor)   | 8.82 | 266.00       |
| ## MTS(ElasticNet)            | 9.64 | 217.62       |

### 3 - 7 Time series cross-validation (cont'd)



### 3 - 7 Time series cross-validation (cont'd)

```
tscv = ns.utils.model_selection.TimeSeriesSplit()


X = np.random.rand(10, 2)

# `fixed_window` = type of slicing
tscv_obj = tscv.split(X, initial_window=3, horizon=2,
                      fixed_window=True)

for train_index, test_index in tscv_obj:
    print("TRAIN:", train_index, "TEST:", test_index)

## TRAIN: [0 1 2] TEST: [3 4]
## TRAIN: [1 2 3] TEST: [4 5]
## TRAIN: [2 3 4] TEST: [5 6]
## TRAIN: [3 4 5] TEST: [6 7]
## TRAIN: [4 5 6] TEST: [7 8]
## TRAIN: [5 6 7] TEST: [8 9]
```

### 3 - 7 Time series cross-validation (cont'd)





### 3 - 7 Time series cross-validation (cont'd)

```
tscv = ns.utils.model_selection.TimeSeriesSplit()

X = np.random.rand(10, 2)

# `fixed_window` = type of slicing
tscv_obj = tscv.split(X, initial_window=3, horizon=2,
                      fixed_window=False)

for train_index, test_index in tscv_obj:
    print("TRAIN:", train_index, "TEST:", test_index)

## TRAIN: [0 1 2] TEST: [3 4]
## TRAIN: [0 1 2 3] TEST: [4 5]
## TRAIN: [0 1 2 3 4] TEST: [5 6]
## TRAIN: [0 1 2 3 4 5] TEST: [6 7]
## TRAIN: [0 1 2 3 4 5 6] TEST: [7 8]
## TRAIN: [0 1 2 3 4 5 6 7] TEST: [8 9]
```